



# NaiveKV: A distributed, persistent key-value store

Yijun Chen, 517021910387, July 2020



上海交通大學

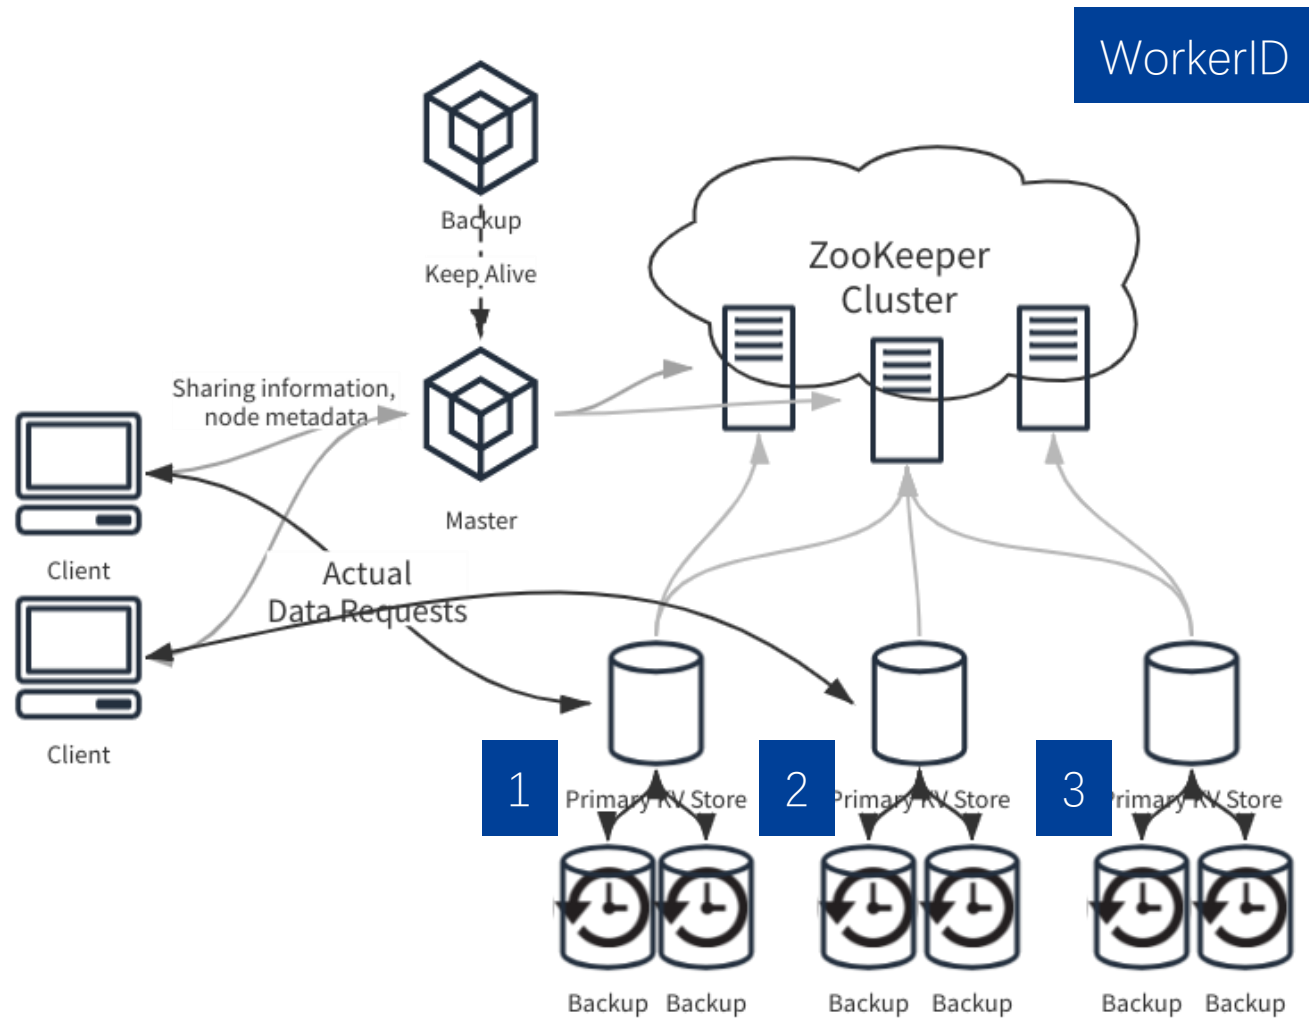
SHANGHAI JIAO TONG UNIVERSITY

- 1 Architecture Overview
- 2 Deploying zookeeper
- 3 Persistence capabilities
- 4 Load balancing
- 5 Migration, backup sync
- 6 Master node high-availability
- 7 Client





# Architecture Overview





# Architecture Overview

- The system is made up of master node(s), primary node(s), backup node(s) and the zookeeper cluster.
- Primary nodes and backup nodes are both “worker” nodes. Primary node and backup node that are in the same group **shares the same Worker ID**, which is **the unit of load balancing and hash slot allocation** for master node.
- Master node(s) and worker node(s) communicate only through modifying metadata stored in zookeeper.



# Deploying ZooKeeper

```
→ zookeeper git:(master) X docker build -t eyek/kv-zookeeper:1.0 .
Sending build context to Docker daemon 9.728kB
Step 1/15 : FROM openjdk:11-jre-slim
---> 973c18dbf567
Step 2/15 : ENV ZOO_CONF_DIR=/conf #...以下省略
```

```
→ zookeeper git:(master) X docker network create zk
→ zookeeper git:(master) X docker run --name zk1 --restart always -d -v $(pwd)/zoo1.cfg:/conf
→ zookeeper git:(master) X docker run --name zk2 --restart always -d -v $(pwd)/zoo2.cfg:/conf
→ zookeeper git:(master) X docker run --name zk3 --restart always -d -v $(pwd)/zoo3.cfg:/conf
→ zookeeper git:(master) X docker container ls
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
1897dd60f7e8	eyek/kv-zookeeper:1.0	"/docker-entrypoint...."	About a minute ago	Up
d8134581bd92	eyek/kv-zookeeper:1.0	"/docker-entrypoint...."	About a minute ago	Up
927ec8157b86	eyek/kv-zookeeper:1.0	"/docker-entrypoint...."	About a minute ago	Up



# Deploying ZooKeeper: configuration

```
clientPort=2181
dataDir=/data
dataLogDir=/data/log
# 默认值2000ms, 为使服务宕机发现更加迅速, 这里设为500ms
tickTime=500
# 以下均为默认配置值
initLimit=5
syncLimit=2
autopurge.snapRetainCount=3
autopurge.purgeInterval=0
maxClientCnxns=60
# 集群中三个节点的地址配置, 其中zoo2与zoo3为Docker分配的域名
server.1=0.0.0.0:2888:3888;2181
server.2=zoo2:2888:3888;2181
server.3=zoo3:2888:3888;2181
```



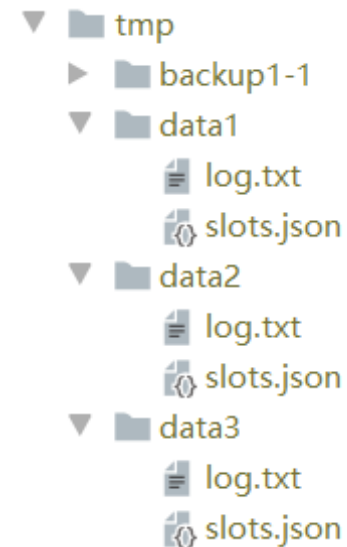
# Metadata in zookeeper

- **/kv**
  - **workers** Registration service, weight information for all workers.
  - **migrations** each round's migration plan
  - **election** playground for backups to re-elect a primary node
  - **masters** Registration service for all masters(and master leader election)
  - **slots** Saves current hash slot allocation table
  - **version** Saves current hash slot allocation table version
  - **workerId** Saves next available workerID



# Persistence capabilities


- To provide fault-tolerance for single node.
- A WAL-based data persistent layer, consists of a base file, a redo log and in-memory lookup cache, like SQLite.
- When committing a change: first change in memory, write log into OS cache, use `sync` to flush change to disk, return.





# Persistence capabilities

- Log format is:
- `<operation> <arg1> <arg2>`  
`<transactionNumber> <versionNumber>`
- Transaction number is here to support concurrent transactions. Will be introduced later.




```
log.txt x
1 put "a" "b" "0" "1"
2 put "c" "d" "0" "2"
3 put "e" "f" "0" "3"
4 put "g" "h" "0" "4"
5 put "i" "j" "0" "5"
6 put "l" "m" "0" "6"
7 put "s" "t" "0" "7"
8 put "y" "z" "0" "8"
```

```
log.txt x
1 start "1"
2 start "2"
3 put "i" "j" "2"
4 commit "2" "1"
5 put "$migration-worker-1" "5" "0" "2"
6 put "o" "p" "1"
7 commit "1" "3"
8 put "$migration-worker-2" "3" "0" "4"
```



# Load Balancing

- Use hash slot allocation. 1024 slots by default.
- Slots allocated according to each worker node's weight.





# Migrating between primary nodes

- Migration happens when new primary node is added into the cluster.
- As described before, every primary node would participate in the migration process.
- Problems:
  - Crash consistency. We need some level of atomicity to ensure no data corruption, and that the migration process can be resumed after either side crashes.
  - Concurrency: It's better not to lock both sides when doing migration, otherwise the whole cluster will be unavailable during migration.

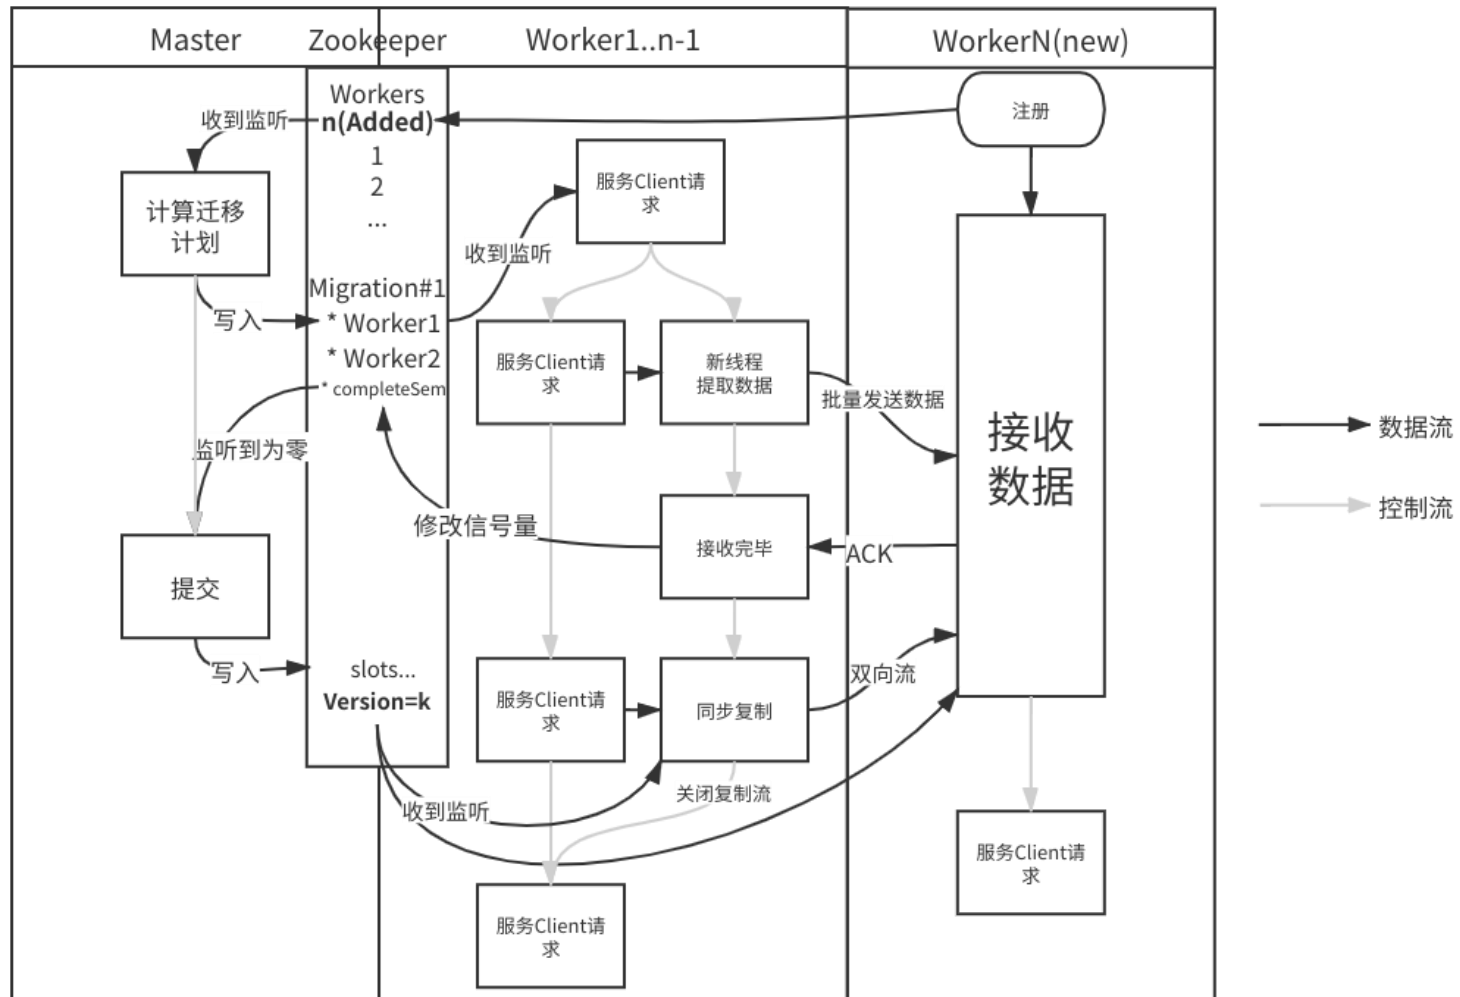


## Two steps of migration

- Full-size replication: Use transaction to maintain all-or-nothing atomicity.
- Incremental replication: Use synchronized replication. Data is first committed to destination worker node before being committed in source worker node.
- Both implemented with gRPC's streaming RPC(long connection).



# Two steps of migration






# Primary-backup synchronization

- Semi-synchronous replication, similar with MySQL.
- If have multiple backups, the primary only have to receive ACK **from one of them** before committing and return to client.
- When primary crash:
  - Backup nodes detect that primary node is down, begin election process.
  - Elect out the one with **highest version number**. Those who have same version number, compare who is first to register in zookeeper.
  - The one got elected becomes primary temporarily, and receives data requests.
  - When the primary is back online, the temporary primary node **sync data with the primary and step down**.



# Primary-backup synchronization

- What if primary and backup crash at the same time?
- No way to guarantee that data is consistent in remaining backup nodes.
- Service downgrade to read only, to avoid undoing.







# Master node high-availability

- Master node provides the client with worker metadata. This part is stateless(zookeeper proxy), easy to replicate.
- Master node also orchestrates the migration process. This require that only one master can operate at the same time.
- Implement a distributed lock in zookeeper's **/kv/masters** directory, with sequential ephemeral zNodes.




# Client, and an optimization

- Client is an REPL which receives user's request from keyboard and send requests to master node and worker node.
- Before: Requesting worker metadata from master before every single requests.
- After: Request all worker metadata from master in a single request. **Include a slot allocation version number in every request sent to worker.** Worker checks the version number, and refuses to serve if version number is wrong. Client refresh its metadata.
- This optimization is implemented after DDL.



# Client, and an optimization



# Implementation and deploying



- Implemented in Go with gRPC, roughly 3k~3.5k lines of code.
- Directly run with `go run ...`
- Hostname and port allocation logic is in Makefile.  
Supports hostname and port customization.

```
weight ?= 10
backupNum ?= 1
primary:
    go run cmd/worker/main.go -id ${id} -port $(shell expr ${id} + 7900) -weight ${weight} -path tmp/data${id}

backup:
    go run cmd/worker/main.go -mode backup -id ${id} -port $(shell expr 10 '*' ${id} + ${backupNum} + 7950) -path
```

谢谢！

